

MD. TANVIR SHEIKH

tanvirsheikh0177@gmail.com • 01778849350 • Dhaka, Bangladesh • [GitHub](#) • [LinkedIn](#)

PROFESSIONAL SUMMARY

AI/ML Engineer with hands-on experience building LLM-powered applications: RAG pipelines (both custom-built and LangChain-based), agent orchestration with LangGraph, and tool/function-calling systems via the Groq API. Comfortable across the full stack — embeddings and vector search (ChromaDB), backend APIs (FastAPI), and deployment (Streamlit, Docker) — with a deep-learning foundation in PyTorch. Currently expanding into production evaluation and cloud deployment.

SKILLS

LLM & GenAI : RAG (retrieval-augmented generation), LangChain, LangGraph, agentic tool use, prompt engineering, LLM-as-judge / answer verification, multi-model fallback routing, LLM APIs (Groq), HuggingFace
Vector Search & Embeddings : ChromaDB, HuggingFace sentence embeddings, MMR search, query expansion
ML Frameworks : PyTorch, Scikit-learn, XGBoost, LightGBM, CatBoost
Backend & Data : FastAPI, Pydantic, Pandas, NumPy, SQL (MySQL, PostgreSQL, SQLite, Supabase)
Tools & Deployment : Git, Docker, Streamlit
Languages : Python

PROJECTS

Personal AI Assistant with RAG — LLM Agent + Retrieval — [• GitHub](#) • [Live Demo](#)

- **Built an agentic conversational assistant** using LangGraph, giving the model access to 3 tools (calculator, web search, RAG retrieval) routed via conditional graph edges — real tool-use orchestration, not a single-call chatbot.
- **Built a multi-model fallback chain** across 3 Groq-hosted models with automatic retry on rate limits, plus an LLM-as-judge verification node (structured Pydantic output) that checks whether each answer actually addresses the question before returning it.
- **Multi-tenant RAG**: per-session knowledge bases, multi-format document ingestion (PDF/DOCX/TXT/MD), query expansion, MMR search, ChromaDB, and HuggingFace embeddings.
- **Hand-rolled token-level streaming** and a custom Streamlit UI with multi-thread chat, LLM-generated thread titles, and session checkpointing.

Tech Stack: Python, LangGraph, LangChain, ChromaDB, HuggingFace, Groq (gpt-oss-120b/20b, qwen3.6-27b), Streamlit, Pydantic

Custom RAG Pipeline — Built from Lower-Level Components — [• GitHub](#)

- **Built a RAG system from first principles** instead of high-level LangChain abstractions — custom EmbeddingManager, VectorStoreManager, and RAGRetriever classes — to demonstrate a working understanding of retrieval internals, not just framework calls.
- **Implemented the full ingestion-to-generation pipeline** (PyMuPDF ingestion, recursive chunking, ChromaDB persistent storage, Llama 3.3 70B via Groq) and debugged production-grade retrieval bugs, including a ChromaDB distance-metric inversion and a silent fallback to default embeddings.

Tech Stack: Python, ChromaDB, SentenceTransformers, PyMuPDF, Groq (Llama 3.3 70B)

Attendance by Selfie — CNN-Based Attendance Platform — [• GitHub](#) • [Live Demo](#)

- **Built an end-to-end system** enabling attendance capture via classroom photo or voice recording, with facial recognition (dlib), speaker verification (Resemblyzer), and a Supabase (PostgreSQL) backend.
- **Designed the database schema and resolved data-integrity issues** (foreign-key errors, query bugs) to keep the Supabase backend reliable across role-based student/teacher portals.

Tech Stack: Python, dlib, Resemblyzer, Supabase, Bcrypt, Segno

F1 Pit Stop Prediction — Kaggle Competition — [• Kaggle](#)

- **Built a stacked ensemble** (LightGBM, XGBoost, CatBoost, PyTorch ANN) with GroupKFold cross-validation and domain-specific feature engineering to improve predictive accuracy.

Tech Stack: Python, LightGBM, XGBoost, CatBoost, PyTorch, GroupKFold

EDUCATION

BSc in Computer Science and Engineering — Alhaj Mockbul Hossain University College, Dhaka

April 2022 – Present

Relevant Coursework: Database Management Systems (DBMS), Data Structures and Algorithms, Artificial Intelligence, Machine Learning, Software Development Engineering.